

CoreDNS

What it does: Provides DNS service inside the Kubernetes cluster so Pods can find each other by name. CoreDNS is the internal DNS of Kubernetes. It allows **Pods and Services to find and talk to each other using names instead of IPs**, which constantly change.

Example: frontend Pod connects to backend.default.svc.cluster.local without knowing IPs.

Real-World Use Case: Microservices Communication

Scenario

You have an **E-commerce application** running on Kubernetes:

- frontend (React)
- backend (Node.js / Java API)
- database (MongoDB)

Each component runs in different Pods, and Pods can restart at any time → **IP addresses change**.

Without CoreDNS (problem)

- Backend Pod IP changes after restart
- Frontend still points to old IP
- Application breaks 

With CoreDNS (solution)

You create a Kubernetes **Service** for backend:

```
apiVersion: v1
kind: Service
metadata:
  name: backend
spec:
  selector:
    app: backend
```

```
  ports:  
    - port: 3000
```

CoreDNS automatically creates a DNS entry:

```
backend.default.svc.cluster.local
```

How traffic flows

1. Frontend sends request to:

```
http://backend:3000
```

2. CoreDNS resolves `backend` → Service IP

3. Service load-balances traffic to backend Pods

4. Application works even if Pods restart 

Another Real-World Example: Database Access

Backend connects to database using DNS:

```
DB_HOST=mongodb.default.svc.cluster.local
```

Even if MongoDB Pods restart or scale, no config change needed.

Enterprise Use Case

- Used in microservices, EKS, GKE, AKS
- Mandatory for service discovery
- Critical for zero-downtime deployments
- Required for HPA, Ingress, Gateway API

Interview One-Liner

CoreDNS enables service discovery in Kubernetes by resolving service names to IPs, allowing microservices to communicate reliably even when Pods change.

storage-provisioner

What it does: Automatically creates **Persistent Volumes (PV)** when a PVC is requested (Dynamic Provisioning). A **storage-provisioner** enables **dynamic storage provisioning** in Kubernetes. It automatically creates **Persistent Volumes (PV)** when an application requests storage using a **PersistentVolumeClaim (PVC)**.

No manual disk creation. No manual PV YAML. Fully automated. 

Example: PVC → automatically creates an EBS volume in AWS or local disk in Minikube.

Real-World Use Case: Database in Kubernetes

Scenario

You are deploying MySQL / MongoDB / PostgreSQL on Kubernetes.

Databases need **persistent storage**, even if Pods restart or move.

Without storage-provisioner (manual & risky)

1. Admin manually creates disk (EBS, NFS, local disk)
2. Admin creates PV YAML
3. App creates PVC
4. Errors if sizes or access modes mismatch

 Slow, error-prone, not scalable

With storage-provisioner (production way)

Step 1: StorageClass (defined once)

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: fast-storage
provisioner: kubernetes.io/aws-ebs
parameters:
  type: gp3
```

Step 2: Application requests storage (PVC)

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mysql-pvc
spec:
  storageClassName: fast-storage
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 20Gi
```

Step 3: What happens automatically

1. PVC created
2. storage-provisioner creates a disk (EBS/NFS/local)
3. PV created and bound
4. Pod mounts storage
5. Data persists even if Pod restarts

📌 Minikube / KIND Example

In Minikube, storage-provisioner creates **local disk storage** automatically.

```
kubectl get pvc
kubectl get pv
```

You'll see PV created without writing PV YAML.

Enterprise Use Cases

- Databases (MySQL, PostgreSQL, MongoDB)
- StatefulSets
- CI/CD tools (Jenkins, Nexus, SonarQube)
- Logging systems (ELK, Prometheus)

ngf-nginx-gateway-fabric

What it does: An **NGINX-based Gateway API implementation** for managing traffic using Kubernetes Gateway API. **ngf-nginx-gateway-fabric** is an **NGINX-based implementation of the Kubernetes Gateway API**. It provides **modern, scalable traffic management** (HTTP/HTTPS routing, TLS, multi-service routing) and is the successor to traditional Ingress.

Example: Routes HTTP traffic to multiple services using `Gateway` and `HTTPRoute` instead of Ingress.

ngf-nginx-gateway-fabric – Real-World Use Case with Example

Think of it as:

Ingress → Next generation traffic control using Gateway API + NGINX

Real-World Use Case: Microservices Traffic Routing

Scenario

You run multiple services in Kubernetes:

- frontend
- catalog
- orders
- payments

All services must be accessed via **one domain** with **path-based routing** and **TLS**.

Old Way (Ingress problem)

- Complex YAML
- Limited extensibility
- Hard to manage multi-team ownership
- Poor separation of infra vs app routing

New Way with NGF (Gateway API)

Step 1: Gateway (Infrastructure team)

```
apiVersion: gateway.networking.k8s.io/v1
kind: Gateway
metadata:
  name: web-gateway
spec:
  gatewayClassName: nginx
  listeners:
  - name: http
    protocol: HTTP
    port: 80
```

Step 2: HTTPRoute (Application team)

```
apiVersion: gateway.networking.k8s.io/v1
kind: HTTPRoute
metadata:
  name: app-routes
spec:
  parentRefs:
  - name: web-gateway
  rules:
  - matches:
    - path:
      type: PathPrefix
      value: /orders
  backendRefs:
```

- name: orders
- port: 8080

Traffic Flow

1. User hits `https://shop.example.com/orders`
2. **NGF (NGINX Gateway Fabric)** receives request
3. Routes traffic using Gateway API rules
4. Forwards request to `orders` service Pod
5. Supports retries, timeouts, TLS, and load balancing 

Enterprise Use Cases

- Multi-tenant Kubernetes clusters
- Platform teams managing gateways
- Microservices traffic routing
- Zero-downtime deployments
- Blue-Green / Canary releases

Why Companies Prefer NGF

- Built on NGINX (battle-tested)
- Uses **Gateway API** (future-proof)
- Clear separation:
 - Infra team → Gateway
 - App team → Routes
- Better than classic Ingress

Interview One-Liner

`ngf-nginx-gateway-fabric` is an NGINX-based Gateway API implementation that provides advanced, scalable traffic management in Kubernetes, replacing traditional Ingress.

Popular Kubernetes Gateway API Implementations

1. Istio Gateway

Use case: Advanced traffic management & service mesh

- Supports **HTTP, HTTPS, TCP**
- Deep integration with **Istio service mesh**
- Ideal for **zero-trust, mTLS, canary, traffic splitting**

Example: Route 90% traffic to v1, 10% to v2 (canary release)

2. Kong Gateway

Use case: API management & external traffic

- Built-in rate **limiting, auth, plugins**
- Works well for **API-first platforms**

Example: Apply OAuth2 + rate limiting on `/api/*` routes

3. Traefik Gateway

Use case: Simple, cloud-native ingress & gateway

- Auto-discovers services
- Lightweight and easy to operate

Example: Expose services automatically when `HTTPRoute` is created

4. Envoy Gateway

Use case: High-performance, cloud-native gateway

- Backed by Envoy Proxy
- Used heavily in cloud providers

Example: Layer 7 routing with retries, timeouts, and observability

5. HAProxy Gateway

Use case: High-throughput, low-latency traffic

- Enterprise-grade reliability
- Popular in finance and telecom

Example: TCP and HTTP routing for legacy + modern apps

6. GKE Gateway (Google Cloud)

Use case: Managed Gateway API on GKE

- Native integration with Google Load Balancers
- Automatic TLS and scaling

Example: Expose services with Google-managed HTTPS LB

7. AWS Gateway API (ALB Controller)

Use case: AWS-native traffic routing

- Uses Application Load Balancer (ALB)
- Integrated with IAM, ACM, WAF

Example: Public HTTPS access via ALB using Gateway API

📌 How to choose (interview-ready)

Need	Best Choice
Simple & NGINX based	NGINX Gateway Fabric
API management	Kong
Service mesh	Istio
High performance	Envoy
Cloud-native (AWS/GCP)	AWS / GKE Gateway

🎯 Interview One-Liner

Gateway API is implemented by multiple providers like NGINX, Istio, Kong, Traefik, Envoy, and cloud-native gateways, each optimized for traffic management, security, and scalability.

kubelet

What it does: Runs on every worker node and ensures Pods and containers are running as defined.

kubelet is the **node-level agent** that runs on **every worker node** in a Kubernetes cluster. It makes sure that **containers described in Pod specs** are actually **running** on that node.

Example: Starts containers, reports Pod status to API Server, restarts failed containers.

Think of kubelet as:

“The executor of Kubernetes instructions on a node.”

🔧 What kubelet is responsible for

- Talks to **kube-apiserver**
- Pulls container images

- Starts & stops containers via container runtime (containerd/CRI-O)
- Performs **health checks** (liveness/readiness)
- Reports Pod & node status
- Mounts volumes and secrets

Real-World Use Case: Application Pod Lifecycle

Scenario

You deploy a backend application:

```
kubectl apply -f backend-deployment.yaml
```

What happens behind the scenes

1. Scheduler assigns the Pod to worker-node-1
2. kubelet on worker-node-1:
 - Pulls the Docker image
 - Creates containers
 - Mounts PVCs & secrets
 - Starts the Pod
3. kubelet continuously checks container health
4. If container crashes → kubelet restarts it automatically 

Example: Health Check

```
livenessProbe:  
  httpGet:  
    path: /health  
    port: 8080
```

If /health fails:

- kubelet kills the container
- kubelet restarts it

Enterprise Use Cases

- Ensures **self-healing applications**
- Enforces resource limits (CPU/memory)
- Enables **autoscaling and resilience**
- Required for StatefulSets, Jobs, and DaemonSets

Where kubelet runs

- Runs as a **system service** on each node
- NOT a Pod
- Communicates securely with API Server

Interview One-Liner

kubelet is the node agent that ensures Pods and containers are running as defined and continuously reports their health and status to the Kubernetes control plane.

kube-proxy

What it does: Handles **network routing and load balancing** for Kubernetes Services.

kube-proxy runs on every worker node and is responsible for **network routing and load balancing** for Kubernetes Services.

Example: A request to `ClusterIP` is forwarded to one of the backend Pods.

Think of kube-proxy as:

“The traffic controller that sends Service requests to the right Pod.”



What kube-proxy is responsible for

- Implements Service networking
- Handles ClusterIP, NodePort, and LoadBalancer traffic
- Load-balances traffic across Pods
- Programs iptables / IPVS rules on nodes
- Enables Pod-to-Service communication



Real-World Use Case: Service Load Balancing

Scenario

You deploy a backend app with 3 replicas:

```
apiVersion: v1
kind: Service
metadata:
  name: backend
spec:
  selector:
    app: backend
  ports:
  - port: 80
```

Pods:

- backend-pod-1
- backend-pod-2
- backend-pod-3



How traffic flows

1. Frontend sends request to:

`http://backend`

2. CoreDNS resolves `backend` → ClusterIP

3. kube-proxy routes traffic to one of the backend Pods

4. Traffic is load-balanced automatically 

NodePort Example

If Service type is NodePort :

```
curl http://<node-ip>:30080
```

kube-proxy forwards the request to a backend Pod.

Enterprise Use Cases

- Microservices communication
- Internal service discovery
- Load balancing without external LB
- Works with Ingress & Gateway APIs

How kube-proxy works internally

Modes:

- **iptables** (default, simple)
- **IPVS** (high performance, scalable)

Where kube-proxy runs

- Runs as a **DaemonSet**
- One Pod per node
- Requires kernel networking support

Interview One-Liner

kube-proxy implements Kubernetes Service networking by routing and load-balancing traffic to Pods using iptables or IPVS rules.

Popular CNI (Container Network Interface) Plugins

(CNI = Pod networking: IP assignment & Pod-to-Pod communication)

◆ Common CNIs

- **Calico** – Networking + Network Policies (widely used in production)
- **Flannel** – Simple overlay networking (beginner-friendly)
- **Cilium** – eBPF-based, high performance, can replace kube-proxy
- **Weave Net** – Simple mesh networking
- **Kindnet** – Default CNI for KIND clusters
- **AWS VPC CNI** – Native AWS networking for EKS
- **Azure CNI** – Native networking for AKS
- **GKE Dataplane v2** – Google's CNI (eBPF-based)
- **Antrea** – VMware-backed, based on Open vSwitch

◆ What CNIs handle

- Pod IP assignment
- Pod-to-Pod communication
- Cross-node networking
- (Some CNIs) Network policies

kube-proxy (Service Networking)

(kube-proxy = Service IP routing & load balancing)

◆ kube-proxy Modes

- **iptables mode** (default)
 - Simple, widely used

- Uses Linux iptables rules
- IPVS mode
 - High performance
 - Better for large clusters
- userspace mode (deprecated ✗)

◆ What kube-proxy handles

- ClusterIP, NodePort, LoadBalancer Services
- L4 load balancing
- Routing Service traffic to Pods

🧠 Special Case: kube-proxy Replacement

Some CNIs replace kube-proxy functionality internally:

- Cilium (eBPF)
- GKE Dataplane v2

→ kube-proxy Pod may not run, but **Service routing still exists**

🎯 Interview One-Liner

CNI plugins provide Pod networking, while kube-proxy (or its replacement) provides Service networking and load balancing in Kubernetes.

🔑 Quick Memory Trick

CNI → Pod IP & Pod traffic
kube-proxy → Service IP & load balancing

kube-scheduler

What it does:

Decides **which node a Pod should run on** based on resources and constraints.

kube-scheduler is the **control plane component** that decides **which worker node a Pod should run on**.

Example:

Schedules a Pod on the node with enough CPU and memory.

Think of kube-scheduler as:

“The brain that places Pods on the right nodes.”

💬 What kube-scheduler considers

When scheduling a Pod, it evaluates:

- * Available **CPU & memory**
- * **Node labels & selectors**
- * **Taints and tolerations**
- * **Affinity / anti-affinity rules**
- * **Topology spread constraints**
- * Resource requests & limits

✅ Real-World Use Case: Placing Pods Correctly

Scenario

You have 3 nodes:

- * `node-1` → high CPU
- * `node-2` → GPU enabled
- * `node-3` → general purpose

A Pod requires a GPU.

```
```yaml
nodeSelector:
 accelerator: nvidia
```



## What happens

1. Pod is created
2. kube-scheduler filters nodes
3. Only node-2 matches
4. Pod is scheduled to node-2 



## Example: High Availability (Anti-Affinity)

---

```
podAntiAffinity:
 requiredDuringSchedulingIgnoredDuringExecution:
 - labelSelector:
 matchLabels:
 app: backend
 topologyKey: kubernetes.io/hostname
```

- ➡ Scheduler spreads Pods across different nodes to avoid single-node failure.



## Enterprise Use Cases

---

- High availability applications
- Cost-optimized workloads
- GPU / special hardware scheduling
- Multi-tenant clusters



## Where kube-scheduler runs

---

- Runs as a **static Pod** on control-plane nodes
- Communicates via **kube-apiserver**
- Does NOT run on worker nodes



## Interview One-Liner

---

kube-scheduler assigns Pods to the most suitable worker nodes based on resource availability, constraints, and policies.

## kube-controller-manager

What it does: Runs controllers that **maintain the desired state** of the cluster.

**kube-controller-manager** runs a set of **controllers** that continuously **ensure the actual cluster state matches the desired state** defined in Kubernetes objects.

Example: If a Pod crashes, the ReplicaSet controller creates a new Pod automatically.

Think of it as:

“The auto-fix engine of Kubernetes.”

## What controllers do

Each controller watches the API Server and **takes action when something is missing or broken**.

Common controllers include:

- Node Controller
- ReplicaSet Controller
- Deployment Controller
- Job Controller
- Endpoint Controller
- Namespace Controller

## Real-World Use Case: Self-Healing Pods

### Scenario

You deploy an app with 3 replicas:

```
spec:
 replicas: 3
```

## What happens

1. One Pod crashes
2. ReplicaSet count becomes 2
3. **ReplicaSet controller** detects mismatch
4. New Pod is created automatically
5. Desired state restored to 3 Pods 

## Another Example: Node Failure

---

- Worker node goes down
- **Node controller** marks it NotReady
- Pods are rescheduled on healthy nodes

## Enterprise Use Cases

---

- Auto-healing applications
- Auto-scaling support
- Continuous reconciliation
- Enforcing cluster policies

## Where kube-controller-manager runs

---

- Runs as a **static Pod** on control-plane nodes
- Communicates via **kube-apiserver**
- Does NOT run on worker nodes

## Interview One-Liner

---

kube-controller-manager continuously monitors the cluster and runs controllers that reconcile the desired state with the actual state.

## Mental Model

---

```
Desired State (YAML)
 ↓
Controller Manager
 ↓
Actual State Fixed Automatically
```

## kube-apiserver

**What it does:** The main entry point of the Kubernetes cluster; all requests go through it. **kube-apiserver** is the front door and brain interface of Kubernetes. Every action in a Kubernetes cluster **must go through the API Server**.

**Example:** `kubectl get pods` → API Server → etcd → response returned.

## kube-apiserver – Real-World Explanation with Example

Think of it as:

“The control center that validates, processes, and stores all cluster requests.”

## What kube-apiserver is responsible for

---

- Exposes the **Kubernetes API**
- Authenticates and authorizes requests (**AuthN/AuthZ**)
- Validates requests and schemas
- Communicates with **etcd** to store/retrieve cluster state
- Acts as the **only component** that talks directly to etcd

## Scenario

You run:

```
kubectl get pods
```

## What happens behind the scenes

1. `kubectl` sends request to `kube-apiserver`
2. API Server:
  - o Authenticates user (cert/IAM/token)
  - o Authorizes via RBAC
3. Reads Pod data from `etcd`
4. Returns response to `kubectl`

 **No API Server = No cluster access**

## Real-World Example: Creating a Pod

---

```
kubectl apply -f app.yaml
```

Flow:

```
kubectl → kube-apiserver → validation → etcd → controllers → kubelet
```

## Enterprise Use Cases

---

- Central access point for:
  - o `kubectl`
  - o CI/CD pipelines
  - o Controllers & schedulers
- Enforces **security and governance**
- Enables multi-tenant clusters
- Required for autoscaling & GitOps

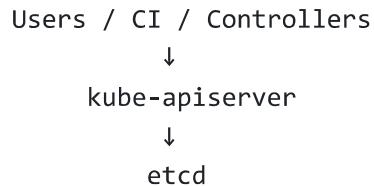
## Where kube-apiserver runs

- Runs as a **static Pod** on control-plane nodes
- Highly available (multiple replicas)
- Exposed securely (TLS)

## Interview One-Liner

kube-apiserver is the central management component of Kubernetes that handles all API requests, enforces security, and stores cluster state in etcd.

## Mental Model



## etcd

**What it does:** A distributed key-value store that stores all cluster state and configuration.

**etcd** is a distributed, consistent key-value database that stores **all** Kubernetes cluster state and configuration.

**Example:** Stores Pod specs, secrets, ConfigMaps, and node information.

Think of it as:

"The single source of truth for the entire Kubernetes cluster."

If etcd is lost → the cluster is lost.

## What is stored in etcd

---

- Pod definitions & status
- Node information
- Deployments, Services, ReplicaSets
- ConfigMaps & Secrets
- RBAC rules
- Cluster configuration & metadata

 Nothing runs without etcd data

## Real-World Use Case: Pod Creation

---

### Scenario

You apply a Deployment:

```
kubectl apply -f app.yaml
```

## What happens

1. Request goes to **kube-apiserver**
2. API Server validates the request
3. **Object is written to etcd**
4. Controllers read from etcd
5. Scheduler schedules Pods
6. kubelet runs Pods

 etcd makes the desired state durable and recoverable 

## Real-World Use Case: Cluster Recovery

---

- Control plane crashes
- New control plane is started

- Reads state from etcd
- Cluster is restored exactly as before

## 🚀 Enterprise Best Practices

---

- Run etcd in **odd numbers** (3 or 5 nodes)
- Enable **TLS encryption**
- Take **regular etcd backups**
- Never expose etcd publicly
- Monitor disk I/O & latency

## 🧠 Where etcd runs

---

- Runs as a **static Pod** on control-plane nodes
- Communicates only with **kube-apiserver**
- Uses **RAFT consensus** for consistency

## 🎯 Interview One-Liner

---

etcd is a highly available, distributed key-value store that holds the complete state of a Kubernetes cluster.

## 🔑 Mental Model

---

```
Cluster State (Truth)
 ↓
 etcd
 ↑
kubernetes only
```

**What it does:** A CNI (Container Network Interface) plugin used by KIND clusters for Pod networking.

**Example:** Assigns IPs to Pods and enables Pod-to-Pod communication.

## metrics-server

**What it does:** Collects CPU and memory usage metrics from nodes and Pods.

**metrics-server** is a cluster add-on that collects CPU and memory usage metrics from Kubernetes nodes and Pods.

**Example:** Used by `kubectl top pods` and Horizontal Pod Autoscaler (HPA).

Think of it as:

“The resource usage meter for Kubernetes.”

Without metrics-server, Kubernetes **cannot autoscale** Pods.

## What metrics-server provides

- CPU usage per Pod
- Memory usage per Pod
- CPU & memory usage per node
- Metrics for HPA (Horizontal Pod Autoscaler)

**!** It does NOT store long-term metrics (that's Prometheus' job).

## Real-World Use Case: Auto Scaling Applications

### Scenario

You run a backend app with HPA:

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
spec:
 scaleTargetRef:
```

```
kind: Deployment
name: backend
minReplicas: 2
maxReplicas: 10
metrics:
- type: Resource
 resource:
 name: cpu
 target:
 type: Utilization
 averageUtilization: 70
```

## What happens

1. metrics-server collects CPU usage from kubelet
2. HPA reads metrics
3. CPU > 70% → Pods scale up
4. CPU < 70% → Pods scale down

 Automatic scaling works only because of metrics-server 

## CLI Example

---

```
kubectl top pods
kubectl top nodes
```

If metrics-server is missing:

```
error: Metrics API not available
```

## Enterprise Use Cases

---

- Horizontal Pod Autoscaling
- Cluster monitoring (basic)
- Resource optimization

- Capacity planning (short-term)

## Where metrics-server runs

- Runs as a **Deployment** in `kube-system`
- Scrapes metrics from **kubelet**
- Uses Metrics API (`metrics.k8s.io`)

## Interview One-Liner

metrics-server collects real-time CPU and memory metrics from Pods and nodes and enables Kubernetes autoscaling.

## Mental Model

`kubelet → metrics-server → HPA / kubectl top`

## Control Plane Components

(Manage and control the Kubernetes cluster)

- **kube-apiserver** – Central API entry point for all cluster operations.
- **etcd** – Distributed key-value store holding cluster state and configuration.
- **kube-scheduler** – Assigns Pods to suitable worker nodes.
- **kube-controller-manager** – Runs controllers to maintain desired cluster state.

## Node (Worker) Components

(Run applications and handle networking on worker nodes)

- **kubelet** – Ensures Pods and containers are running as defined.

- **kube-proxy** – Handles Service networking and load balancing.
- **kindnet (CNI)** – Provides Pod-to-Pod networking and IP assignment.

## Add-ons / Cluster Services

---

(Extend Kubernetes functionality)

- **CoreDNS** – Provides DNS-based service discovery inside the cluster.
- **metrics-server** – Collects CPU and memory metrics for Pods and nodes.
- **storage-provisioner** – Enables dynamic Persistent Volume provisioning.
- **ngf-nginx-gateway-fabric** – Implements Gateway API using NGINX for traffic routing.