


🧠 What is a Container Runtime?



A container runtime is the component responsible for:

- Pulling container images
- Creating containers
- Running containers
- Stopping & deleting containers

👉 Kubernetes does NOT run containers directly 👉 It talks to the runtime using **CRI** (Container Runtime Interface)

E Container Runtime Architecture





🔗 Container Runtime Landscape

Runtime	Type	CRI	Production
Docker	High-level	✗ (removed)	✗
containerd	Low-level	✓	★★★★★
CRI-O	Low-level	✓	★★★★★
runc	OCI runtime	✗	internal

Runtime	Type	CRI	Production
Kata Containers	VM-based	✓	niche

1. Docker (✗ Deprecated in Kubernetes)

! Important Truth

- Docker was removed from Kubernetes v1.24+
- Kubernetes never used Docker directly
- It used `dockershim` (now removed)

Kubernetes ✗ Docker
Kubernetes ✓ containerd / CRI-O

✓ Docker still useful for local development

2. containerd (🔥 DEFAULT & RECOMMENDED)

Why containerd?

- CNCF project
- Lightweight
- High performance
- Used by EKS, GKE, AKS, kubeadm

Install containerd (Ubuntu – Production Way)

Step 1: Install

```
sudo apt update
```

```
sudo apt install -y containerd
```

Step 2: Generate config

```
sudo mkdir -p /etc/containerd  
containerd config default | sudo tee /etc/containerd/config.toml
```

Step 3: Enable Systemd Cgroup (MANDATORY)

Edit:

```
sudo vi /etc/containerd/config.toml
```

Set:

```
SystemdCgroup = true
```

Step 4: Restart

```
sudo systemctl restart containerd  
sudo systemctl enable containerd
```

🔍 Verify containerd

```
crtictl info  
crtictl ps
```

📜 Kubernetes Uses containerd Like This

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: runtime-test  
spec:
```

```
containers:  
- name: nginx  
  image: nginx
```

⬆ Kubernetes pulls & runs image via **containerd + runc**

3. CRI-O (RedHat / OpenShift)

Why CRI-O?

- Built **only** for Kubernetes
- No extra features
- Secure by default

Used by:

- OpenShift
- Some hardened clusters

🔧 Install CRI-O (Ubuntu Example)

```
sudo apt install -y cri-o cri-o-runc  
sudo systemctl enable crio  
sudo systemctl start crio
```

Verify:

```
crictl info
```

📜 Pod YAML (Same as containerd)

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: crio-test
```

```
spec:  
  containers:  
  - name: busybox  
    image: busybox  
    command: ["sleep", "3600"]
```



4. runc (Low-Level Runtime)

What is runc?

- Executes containers

- Implements OCI spec

- Used internally by:

- containerd

- CRI-O

- Docker

You never configure runc directly in Kubernetes.



Runtime Relationship (VERY IMPORTANT)

Docker

```
└ containerd  
  └ runc
```

Kubernetes

```
└ containerd / CRI-O  
  └ runc
```



5. Kata Containers (Extra Isolation)

What makes Kata special?

- Each container runs in a **lightweight VM**
- Strong isolation
- Slight performance overhead

Use cases:

- Multi-tenant clusters
- Untrusted workloads

RuntimeClass Example (Kata)

```
apiVersion: node.k8s.io/v1
kind: RuntimeClass
metadata:
  name: kata
handler: kata-runtime
```

Pod using Kata:

```
spec:
  runtimeClassName: kata
```

How to Check Runtime in Your Cluster

```
kubectl get nodes -o wide
```

Look for:

```
CONTAINER-RUNTIME
containerd://1.7.x
```

Or:

```
kubectl describe node | grep -i runtime
```



Common Runtime Issues (On-Call SRE)

Issue	Cause
Pods stuck in ContainerCreating	Runtime down
CrashLoopBackOff	Image / runtime error
kubelet not starting	Cgroup mismatch
ImagePullBackOff	containerd registry issue

Fix:

```
sudo systemctl restart containerd  
journalctl -u containerd
```



Runtime Selection Guide

Environment	Runtime
kubeadm	containerd
EKS	containerd
GKE	containerd
OpenShift	CRI-O
Multi-tenant	Kata
Local Dev	Docker



What is containerd?


containerd is a **low-level container runtime** responsible for:




- Pulling container images


- Creating containers
- Running containers
- Managing container lifecycle

👉 Kubernetes talks to containerd using CRI (Container Runtime Interface) 👉 containerd internally uses runc to start containers

Container Runtime Architecture







Install containerd (Ubuntu)

Step 1: Install containerd package

```
sudo apt install -y containerd.io
```

Step 2: Create containerd config directory

```
sudo mkdir -p /etc/containerd
```

Step 3: Generate default configuration

```
sudo containerd config default | sudo tee /etc/containerd/config.toml
```

This creates the [official default runtime configuration](#).

Step 4: Restart & enable containerd

```
sudo systemctl restart containerd
sudo systemctl enable containerd
```

Step 5: Verify service

```
systemctl status containerd
```

Pull Image Using `ctr`

```
ctr image pull docker.io/library/alpine:latest
```

Notes

- `ctr` is `containerd`'s native CLI
- Mostly for debugging
- Not user-friendly
- **Not recommended** for daily use

Run Container Using `nerdctl`

```
nerdctl run -it --rm docker.io/library/alpine:latest
```

Why `nerdctl`?

- Docker-like syntax
- Works directly with `containerd`
- Supports:
 - volumes
 - networks
 - compose

Tool	Ease	Use Case
ctr	 Hard	Debugging
nerdctl	 Easy	Daily usage
docker	 Easy	Dev only

CRI Tool – `cricctl`

`cricctl`

What is `cricctl`?

- CLI for CRI-compatible runtimes
- Used by Kubernetes admins
- Talks directly to containerd / CRI-O

Common commands:

```
cricctl info  
cricctl images  
cricctl ps  
cricctl pods
```

Tool Comparison (Very Important)

Tool	Talks To	Used By
ctr	containerd	Runtime debugging
nerdctl	containerd	Humans
cricctl	CRI	Kubernetes admins
kubectl	kube-apiserver	Users

Production Best Practices

- ✓ Use `containerd`, not Docker
- ✓ Use `nerdctl` for local testing
- ✓ Use `cricctl` for Kubernetes debugging
- ✓ Never use `ctr` in automation
- ✓ Always enable `systemd cgroups` in `/etc/containerd/config.toml`



Mandatory Production Setting (IMPORTANT)

Edit:

```
sudo vi /etc/containerd/config.toml
```

Set:

```
SystemdCgroup = true
```

Restart:

```
sudo systemctl restart containerd
```

Without this → kubelet will fail



How Kubernetes Uses containerd

```
apiVersion: v1
kind: Pod
metadata:
  name: runtime-test
spec:
  containers:
  - name: alpine
    image: alpine
    command: ["sleep", "3600"]
```

Flow:

```
kubectl → kubelet → CRI → containerd → runc → kernel
```

Common On-Call Issues

Issue	Cause
ImagePullBackOff	Registry / auth issue
ContainerCreating	Runtime down
kubelet crash	Cgroup mismatch
Pods stuck	containerd not running

Fix:

```
sudo systemctl restart containerd  
journalctl -u containerd
```

Key Takeaways (Interview Gold)

- Kubernetes does NOT use Docker
- containerd is production standard
- runc actually creates containers
- CRI is the runtime interface
- nerdctl = Docker for containerd