



Kubernetes Autoscaling – HPA & VPA (Real-World Guide)

This repository demonstrates **Horizontal Pod Autoscaler (HPA)** and **Vertical Pod Autoscaler (VPA)** with real YAML manifests and production architecture.



Table of Contents

- Autoscaling Overview
- HPA vs VPA (Truth Table)
- Autoscaling Architecture
- Horizontal Pod Autoscaler (HPA)
- Vertical Pod Autoscaler (VPA)
- Installation Prerequisites
- YAML Examples
- Verification & Debugging
- Production Best Practices
- When NOT to use Autoscaling



Autoscaling Overview

Kubernetes supports **two main autoscaling mechanisms**:


Autoscaler	Scales	Direction
HPA	Pods	Horizontal
VPA	CPU / Memory	Vertical



 **⚠️ HPA and VPA should NOT manage the same resource simultaneously**

vs HPA vs VPA (Critical Difference)

Feature	HPA	VPA
Scales pods	✓	✗
Scales CPU/Memory	✗	✓
Restart pods	✗	✓
Uses Metrics Server	✓	✗
Used in production	★★★★★	★★★
Stateful workloads	✗	⚠

Autoscaling Architecture





Metrics Server

Horizontal Pod Autoscaler (HPA)

🔍 What HPA Does

- Automatically increases/decreases **replica count**
- Based on:
 - CPU
 - Memory

- Custom metrics (Prometheus)

HPA Prerequisites

1 Metrics Server (Mandatory)

```
kubectl get deployment metrics-server -n kube-system
```

If missing:

```
kubectl apply -f https://github.com/kubernetes-sigs/metrics-
server/releases/latest/download/components.yaml
```

2 Resource Requests MUST Exist

```
resources:
  requests:
    cpu: "100m"
    memory: "128Mi"
```

Without this → HPA will not work

HPA YAML (hpa.yml)

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: app-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: demo-app
  minReplicas: 2
  maxReplicas: 10
  metrics:
```

```
- type: Resource
  resource:
    name: cpu
    target:
      type: Utilization
      averageUtilization: 60
```

How HPA Works (Flow)

CPU > 60% → scale UP
CPU < 60% → scale DOWN

Verify HPA

```
kubectl get hpa
kubectl describe hpa app-hpa
```

Live watch:

```
kubectl get pods -w
```

Vertical Pod Autoscaler (VPA)

What VPA Does

- Automatically adjusts **CPU & memory**
- Based on historical usage
- Requires **pod restart**



VPA Limitations (Important)

Limitation	Impact
Pod restart	Downtime risk
Not for stateless	Avoid
Not with HPA	Conflict



VPA Installation (One-Time)

```
kubectl apply -f  
https://github.com/kubernetes/autoscaler/releases/latest/download/vertical-pod-  
autoscaler.yaml
```

Verify:

```
kubectl get pods -n kube-system | grep vpa
```



VPA YAML (vpa.yaml)

```
apiVersion: autoscaling.k8s.io/v1
kind: VerticalPodAutoscaler
metadata:
  name: app-vpa
spec:
  targetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: demo-app
  updatePolicy:
    updateMode: Auto
```

VPA Update Modes

Mode	Behavior
Off	Only recommendations
Initial	On pod creation
Auto	Live updates (restart)

Verify VPA

```
kubectl describe vpa app-vpa
```

Check recommendations:

```
kubectl get vpa
```

Real-World Use Cases

Scenario	Use
Web apps	HPA
APIs	HPA
Batch jobs	VPA
ML workloads	VPA
Stateful apps	



Common Autoscaling Issues (On-Call)

Issue	Cause
HPA not scaling	No resource requests
VPA not updating	Mode set to Off
Pods restarting	VPA Auto mode
Metrics unavailable	Metrics server down

Fix:

```
kubectl top pods  
kubectl top nodes
```



Production Best Practices

- ✓ Always set **requests & limits**
- ✓ Prefer HPA for stateless apps
- ✓ Use VPA in recommendation mode first
- ✓ Never combine HPA + VPA on CPU
- ✓ Use Cluster Autoscaler with HPA
- ✓ Monitor scaling events



Autoscaling Decision Guide

Requirement	Choose
Handle traffic spikes	HPA
Optimize resource usage	VPA
Zero downtime	HPA
Cost optimization	VPA
High availability	HPA



Interview-Ready One-Liners

- HPA scales **pods**
- VPA scales **resources**
- HPA needs **metrics-server**
- VPA restarts pods
- Never mix HPA & VPA on CPU