

# Kubernetes Resources – Requests & Limits (Deep Dive)

---

This document explains how **CPU and Memory Requests & Limits** work in Kubernetes, how the **scheduler** and **kubelet** enforce them, and what happens during **OOMKill** and **throttling**.

## Table of Contents

---

- What Are Requests & Limits?
- Scheduler vs Kubelet Responsibilities
- Architecture Diagram
- Requests vs Limits (Truth Table)
- YAML Examples
- Memory OOM Scenario
- CPU Throttling Scenario
- Verification & Debugging
- Production Best Practices
- Common On-Call Incidents

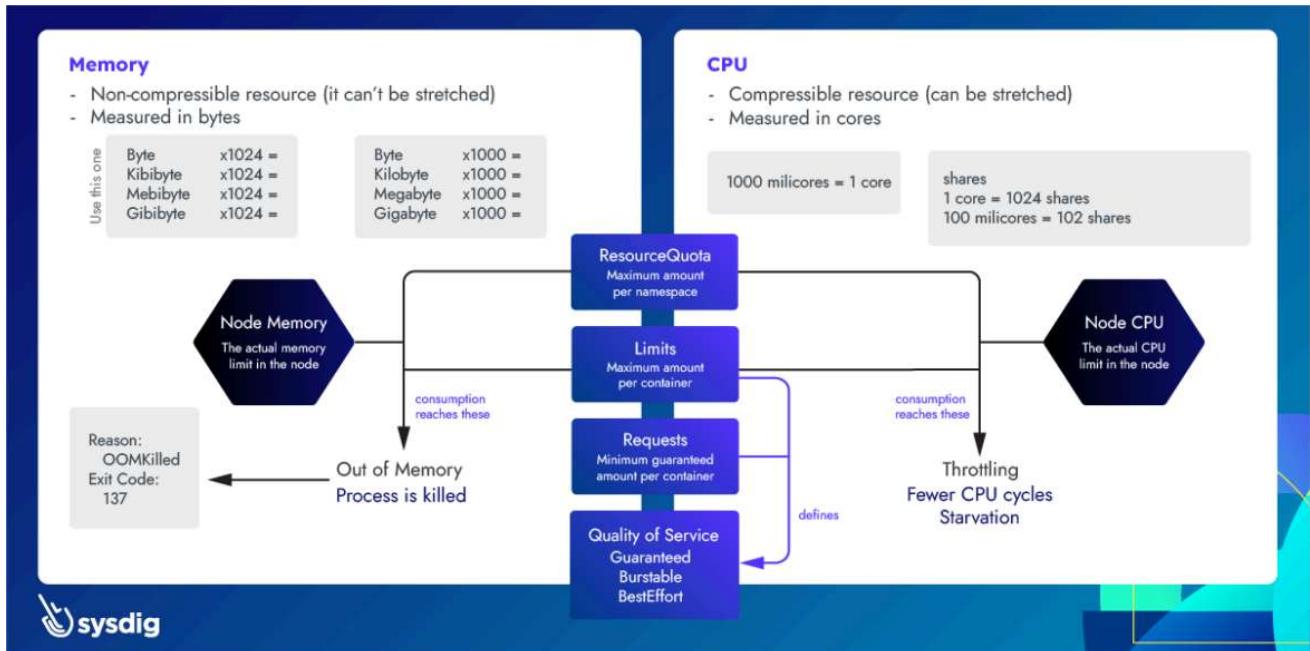
## What Are Requests & Limits?

---

Resource	Request	Limit
CPU	Guaranteed minimum	Maximum allowed
Memory	Guaranteed minimum	Hard upper bound

Request = Scheduling decision Limit = Runtime enforcement

# Resource Management Architecture



The pod - `Deployment.yaml`

```

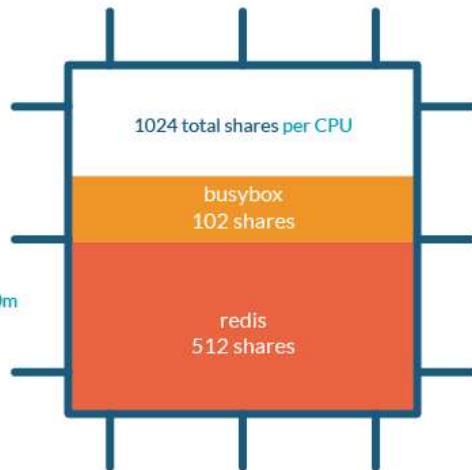
kind: Deployment
apiVersion: extensions/v1beta1
metadata:
  name: redis
...
  template:
    spec:
      containers:
        - name: redis
          image: redis:5.0.3-alpine
          resources:
            requests:
              memory: 300Mi
              cpu: 500m
...
        - name: busybox
          image: busybox:1.28
          resources:
            requests:
              memory: 100Mi
              cpu: 100m
  
```

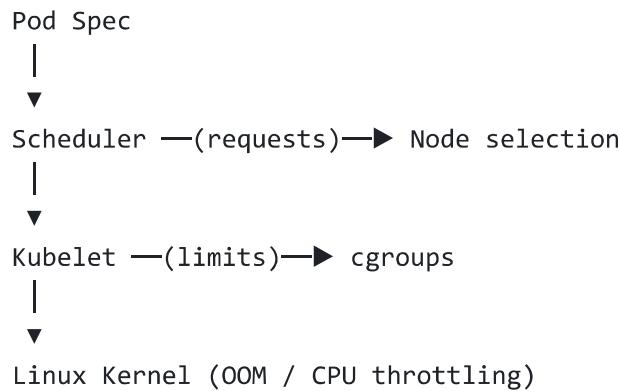
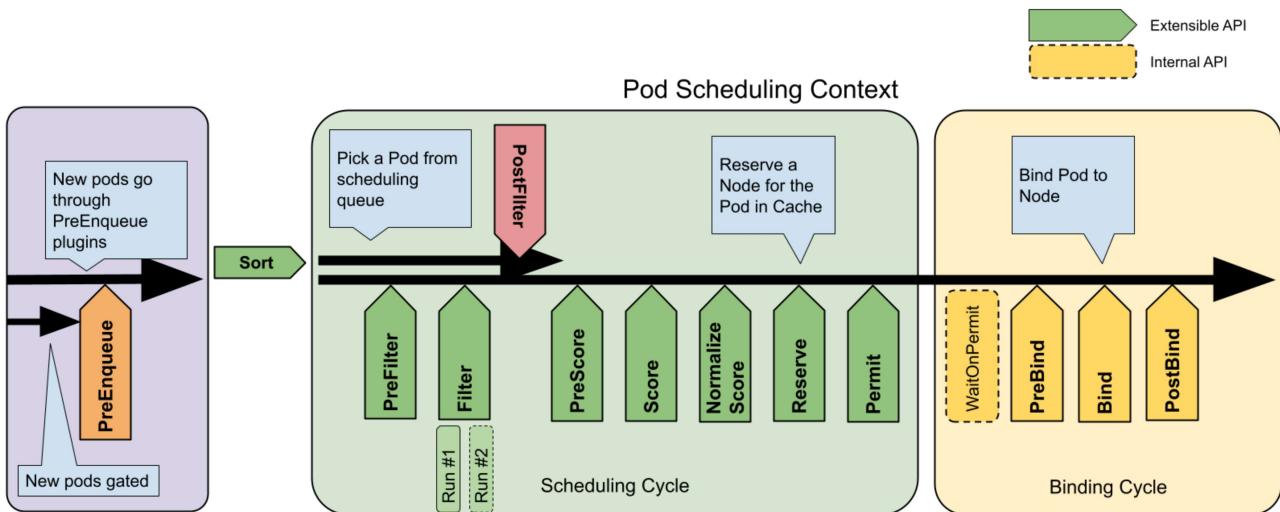
## CPU Requests

Kubernetes assigns 1024 shares per core.  
 $1 \text{ core} = 1000 \text{ milicores} = 1000m$

$1024 * 0.5 = 512 \text{ shares}$

$1024 * 0.1 = 102 \text{ shares}$





## ⚖️ Requests vs Limits (Critical Truth Table)

Scenario	Result
Request > Node capacity	Pod Pending
Memory > Limit	OOMKilled
CPU > Limit	Throttled
No requests	BestEffort Pod
No limits	Unlimited usage



# YAML Example – Requests & Limits ( request-limits.yml )

---

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: stress-test
  labels:
    tier: frontend
spec:
  replicas: 3
  selector:
    matchLabels:
      app: stress-test
  template:
    metadata:
      labels:
        app: stress-test
        tier: frontend
    spec:
      containers:
        - name: stress-test
          image: polinux/stress
          resources:
            requests:
              memory: "64Mi"
              cpu: "250m"
            limits:
              memory: "128Mi"
              cpu: "500m"
          command:
            ["stress", "--vm", "1", "--vm-bytes", "256M", "--timeout", "60s"]
```

## Meaning:

- Pod **guaranteed** 200m CPU, 256Mi memory
- Pod **cannot exceed** 500m CPU, 512Mi memory



## Scheduler Behavior (IMPORTANT)

---

Scheduler only looks at **requests**, never limits.

```
Node Capacity = 4 CPU / 8Gi
Pod Requests = 500m CPU / 1Gi
```

- Scheduler allows 8 such pods → Limits are ignored during scheduling

## 💥 Memory Over Limit – OOM Scenario

(`request-limits-over-mem-limits.yml`)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: stress-test
  labels:
    tier: frontend
spec:
  replicas: 3
  selector:
    matchLabels:
      app: stress-test
  template:
    metadata:
      labels:
        app: stress-test
        tier: frontend
    spec:
      containers:
        - name: stress-test
          image: polinux/stress
          resources:
            requests:
              memory: "3000Mi"
              cpu: "250m"
            limits:
              memory: "3000Mi"
              cpu: "500m"
          command:
            ["stress", "--vm", "1", "--vm-bytes", "4096M", "--timeout", "60s"]
```

What happens?

- App crosses 256Mi

- Kernel triggers **OOMKill**
- Pod restarts
- Status:

OOMKilled

Check:

```
kubectl describe pod <pod-name>
```

## Memory Is NOT Throttled

---

CPU → throttled  
Memory → killed

There is **NO** memory throttling in Linux

## CPU Limit – Throttling Scenario

---

```
resources:  
  limits:  
    cpu: "500m"
```

When app uses more:

- Pod is slowed down
- No restart
- No error logs
- Hard to detect

Check throttling:

```
kubectl describe pod
```

Or Prometheus:

```
container_cpu_cfs_throttled_seconds_total
```

## 🎯 QoS Classes (VERY IMPORTANT)

---

QoS Class	Condition
Guaranteed	Requests = Limits
Burstable	Requests < Limits
BestEffort	No requests/limits

```
kubectl get pod <pod> -o jsonpath='{.status.qosClass}'
```

## ⚠️ Real On-Call Scenarios

---

Issue	Root Cause
Pod Pending	Requests too high
Random restarts	Memory limit too low
Slow app	CPU throttling
Node OOM	No memory limits
HPA not working	Missing requests

## 🔍 Debugging Commands (SRE)

---

```
kubectl top pods  
kubectl top nodes  
kubectl describe pod <pod>  
kubectl get events
```

Check OOM:

```
kubectl logs <pod> --previous
```

## ✓ Production Best Practices

- ✓ Always set **requests**
- ✓ Always set **memory limits**
- ✓ Avoid CPU limits for latency-sensitive apps
- ✓ Use **Burstable** QoS for web apps
- ✓ Use **Guaranteed** QoS for critical services
- ✓ Use **VPA recommendations**
- ✓ Monitor OOMKills aggressively

## 🧭 Recommended Values (Real-World)

App Type	CPU Request	CPU Limit	Memory
Web API	200m	✗	512Mi
Batch Job	500m	1	1Gi
Database	1	1	4Gi
JVM App	500m	✗	2Gi

✗ CPU limits often removed in production

## 🧠 Interview-Ready One-Liners

- Scheduler only cares about **requests**
- Memory limit breach = **OOMKill**
- CPU limit breach = **throttling**
- BestEffort pods die first
- Guaranteed pods die last